



601.220 Intermediate Programming

Summer 2022, Meeting 7 (June 22nd)

Today's agenda

- Exercises 11 and 12 review
- “Day 13” material
 - Lifetime/scope, struct types, random number generation
 - Exercise 13
- “Day 14” material
 - Binary file I/O, bitwise operations
 - Exercise 14

Reminders/Announcements

- HW3 due today
- Midterm project teams
 - If you were not registered on a team as of 11pm last night, you will be assigned to a team today
 - Team repositories will be available by end of day today
- Midterm project: overview in class on Friday June 24th, due Friday, July 1st

Exercise 11 review

pairwise_sum.c: When running the program using valgrind:

```
valgrind --leak-check=full ./pairwise_sum
```

A memory leak is reported:

```
==17736== 16 bytes in 1 blocks are definitely lost in loss record 1 of 1
==17736==    at 0x483B7F3: malloc (in /usr/lib/x86_64-linux-gnu/valgrind/vgpreload_memcheck.so.3)
==17736==    by 0x10922B: pairwise_sum (pairwise_sum.c:28)
==17736==    by 0x109399: main (pairwise_sum.c:57)
```

valgrind indicates there is a memory leak: the memory is allocated in

Exercise 11 review

In the code:

```
int *pairsum2 = pairwise_sum(pairwise_sum(array, 5), 4);  
// ...  
free(pairsum2);
```

leak

Issue: `pairwise_sum` returns a pointer to a dynamically allocated array, but for the “inner” call, the array is never freed.

Fix:

```
→ int *a = pairwise_sum(array, 5);  
int *pairsum2 = pairwise_sum(a, 4);  
// ...  
free(pairsum2);  
→ free(a);
```

Exercise 11 review

`primes.c`:


Issue: the `set_primes` function needs to call `realloc` if the array of results needs to be increased in size.

However, `realloc` can and usually does return a pointer to a new dynamic array (with a different memory address).

Unless `set_primes` can modify the list pointer in `main`, the `main` function has no way of knowing the address of the re-allocated array.

Exercise 11 review

Sketch showing the problem with the original code:



Exercise 11 review

Solution: change `set_primes` so that it takes a pointer to the list pointer variable in the main function.

// set_primes function: originally

```
int set_primes( int *list , int capacity )
```

// updated

```
int set_primes( int **list , int capacity )
```

// in main function

```
int *list = /* initial allocation of array */
```

// original call to set_primes

```
int prime_count = set_primes( list , capacity );
```

// updated

```
int prime_count = set_primes( &list , capacity );
```



Exercise 11 review

Sketch showing how having `set_primes` take a pointer to a pointer solves the problem:

set_primes

`list` \rightarrow `*list`

`*list = realloc(...)`



Exercise 11 review

Changes to `set_primes`: essentially, everywhere that `list` was mentioned, we now want `*list` so that we are referring (indirectly) to the `list` pointer variable in `main`.

One issue: array subscript operator has higher precedence than the pointer dereference operator (`*`)

So, instead of changing

```
list[idx++] = n;
```

to

```
*list[idx++] = n;
```

it should be

```
(*list)[idx++] = n;
```

`list` \Rightarrow `(*list)`

Exercise 12 review

Declaration of search function:

How it is called:

```
pos = search(arr1, arr1 + 10, 318);
```

Declaration:

```
int *search(int *start, int *end, int searchval);
```

Exercise 12 review

Useful property when lower bound of search range is inclusive, and upper bound is exclusive: end - start is the number of elements in the range. So:

```
int *search(int *start, int *end, int searchval) {
    int num_elts = (int) (end - start);
    if (num_elts < 1) {
        return NULL; // no elements in range
    } else {
        // general case: check middle element, if it's equal to
        // searchval, success, otherwise continue recursively on
        // left or right side of range
    }
}
```

Exercise 12 review

```
// search, general case
int *mid = start + (num_elts/2);
if (*mid == searchval) {
    return mid; // success, found the search value
} else if (*mid < searchval) {
    // continue recursively in right side of range
} else {
    // continue recursively in left side of range
}
```

Exercise 12 review

equivalence

$$a[i] \equiv *(a + i)$$



```
// in the test code, finding the index of the matching element
pos = search(arr1, arr1 + 10, 318);
assert(pos != NULL);
assert(*pos == 318);
// TODO: compute the index of the matching element
index = pos - arr1; // <-- add this
assert(2 == index);
```

Exercise 12 review

General observation about 2-D arrays: if p is a pointer to an element, and N is the number of columns in one row, then

$p + N$

yields a pointer to an element that is in the same column and next row from the element p points to. Picture:



"row major"

`int a[3][5];`
 $\dots \rightarrow a[i]$

Exercise 12 review

makeCol:

```
// TODO: declare the unit variable (array of 9 integers, to be returned)  
int *unit = malloc(9 * sizeof(int));
```


Exercise 12 review

makeCube:

```
// TODO: declare the unit variable (array of 9 integers, to be returned  
int *unit = malloc(9 * sizeof(int));
```



Exercise 12 review

checkRows:

```
// TODO: call check on current row and add to variable good  
good += check(&table[r][0]);
```

Observation: ^{table[r]}elements in a single row are contiguous in memory
(each row of a 2-D array can be treated as a 1-D array).

Exercise 12 review

checkCols:

```
for (int c = 0; c < SIZE; c++) {  
    // TODO: call makeCol on current column and assign result to column  
    column = makeCol(&table[0][c]);    // <-- get one column of values  
    good += check(column);  
    free(column);    // <-- free dynamic array  
}
```

Exercise 12 review

checkCubes:

```
// TODO: call makeCube on current cube and assign result to variable cube  
cube = makeCube(&table[r][c]);    // <-- get 3x3 "cube" of values  
good += check(cube);  
free(cube);                       // <-- free dynamic array
```

Exercise 12 review

main (in sudoku.c): code does not call `fclose` to close input file: should modify main function so that `infile` is guaranteed to be closed (using `fclose`) if it is opened successfully.

Makefile: `CFLAGS` should include the `-g` option (to enable debug symbols).

Running `valgrind`:

```
valgrind ./main --leak-check=full --show-leak-kinds=all <name of input file>
```

Day 13 recap questions

- ① What is *struct* in C?
- ② How are the fields of a struct passed into a function - by value or by reference?
- ③ What is the size of a *struct*? What is structure padding in C?
- ④ What is the difference between lifetime and scope of a variable?
- ⑤ What is variable shadowing (i.e. hiding)?
- ⑥ What is the output of the below program?

1. What is *struct* in C?

struct introduces a *used-defined data type*.

Very much like a class in Java or Python, but with only the ability to include member variables, not member functions.


An instance of a struct is a “bundle” of variables that are packaged as a single entity.

Example:

```
struct Point {  
    int x, y;  
};
```

Initialization

```
// ... elsewhere in the program ...  
struct Point p = { .x = 2, .y = 3 };
```




2. How are the fields of a struct passed into a function - by value or by reference?

Instances of a struct type are passed by value. E.g.

```
struct Point { int x, y; };
```

```
void f(struct Point p, int dx) {  
    p.x += dx;  
}
```

```
int main(void) {  
    struct Point q = { .x = 4, .y = 5 };  
    f(q, -2);  
    printf("%d,%d\n", q.x, q.y); // prints "4,5"  
    return 0;  
}
```



3. What is the size of a *struct*? What is structure padding in C?




`sizeof(struct Foo)` is the sum of the sizes of the fields of struct Foo, plus the total size of any padding inserted by the compiler to ensure that fields are correctly aligned.

alignment: the memory address of a variable (including a field variable in an instance of a struct type) must be a multiple of the size of the field.

E.g., a 4-byte int variable (or struct field) must have its storage allocated starting at a machine address that is a multiple of 4.


The compiler will insert padding automatically: you don't need to do anything special. `sizeof(struct Foo)` will always take the padding into account. Just trust that the compiler will figure out the right struct layout to use.


struct padding example

```
1 struct Foo {  
4   char a;   
8   int b;    
   long c;    
};
```

```
// ...
```

```
struct Foo f;  
printf("%lu\n", sizeof(f));
```





4. What is the difference between lifetime and scope of a variable?

"dynamic"

Lifetime: the interval from (1) the point in time when a variable is created, to (2) the point in time when a variable is destroyed. Examples:

- the lifetime of a local variable is the duration of the function call
- the lifetime of a global variable is the duration of the entire program


"static"

Scope: the region of the program code in which a variable may be accessed. Examples:

- the scope of a local variable is from its declaration to the closing "}" of the block in which it's defined
- the scope of a global variable is the entire program (assuming that there is a declaration or definition of the variable in the current block, or in the enclosing block}

5. What is variable shadowing (i.e. hiding)?

Shadowing: a variable declaration in a nested scope has the same name as a variable in an “outer” scope.



Shadowing example

```
int x;

void foo(int x) {
    {
        int x = 5;
        printf("%d\n", x); // prints "5"
    }
    printf("%d\n", x); // prints "4"
}

int main(void) {
    x = 3;
    foo(4);
    printf("%d\n", x); // prints "3"
    return 0;
}
```

The diagram illustrates variable shadowing in C. It shows a global variable `x` and a local variable `x` inside a function `foo`. Hand-drawn blue annotations include arrows showing the scope of each variable, boxes around the values 5 and 3, and a dashed line indicating the shadowing relationship.

6. What is the output of the below program?

```
#include <stdio.h>
int foo;
void bar() {
    int foo = 3;
    {
        extern int foo;
        printf("%d; ", foo);
        foo = 2;
    }
    printf("%d; ", foo);
}
void baz() { printf("%d; ", foo); }
int main() {
    {
        int foo = 5;
        bar();
        printf("%d; ", foo);
    }
    baz();
    return 0;
}
```

The diagram illustrates the scope resolution for the variable `foo`. A solid blue circle highlights the local declaration `int foo = 3;` inside the `bar()` function. A dashed blue line with an arrow points from this local declaration to the `extern int foo;` statement, indicating that the local variable shadows the global one. Another solid blue line with an arrow points from the `extern int foo;` statement to the `printf("%d; ", foo);` call within the same block, showing that the program uses the global `foo` (which has the value 5 from `main`) instead of the local one.

. VS. ->

To access a member variable of a struct instance directly, use the “.” operator. To access a member variable of a struct instance indirectly via a pointer, use the -> operator.

Note that $p \rightarrow x$ means exactly the same thing as $(*p) . x$. It's just a more convenient syntax.



Example of . vs. ->

```
struct Player { int x; int y; int health; };
```

```
struct Player player;
```

```
player.x = 42;
```

```
player.y = 17;
```

```
struct Player *p = &player;
```

```
p->health = 100;
```


Exercise 13

meet back 11:30

- Working with struct types, including pointers to instances of struct types
- Breakout rooms 1–10 are “social”
- Use Slack to let us know if you have any questions!

Day 14 recap questions

- ① How do we read/write binary files in C?
- ② What character represents the bitwise XOR operation? How does it differ from the OR operation?
- ③ What happens if you apply the bitwise operation on an integer value? (extra: what if we apply to floats)
- ④ What is the result of `(15 >> 2) || 7`?
- ⑤ What is the result of `(15 >> 2) | 7`?

1. How do we read/write binary files in C?

983

Use "rb" or "wb" when calling `fopen`, and use `fread` or `fwrite` to read binary data value(s).

E.g., read array of 20 `int` values from a file of binary data:

```
FILE *in = fopen("input.dat", "rb");
if (in == NULL) { /* handle error */ }
else {
    int *arr = malloc(sizeof(int) * 20);
    int rc = fread(arr, sizeof(int), 20, in);
    if (rc != 20) { /* handle error */ }
}
// arr now points to array of 20 elements read from input file
```

Warning: this code is not portable across CPU architectures due to byte ordering. For multi-byte data values some CPUs store least significant byte first ("little endian"), some CPUs store most significant byte first ("big endian")

Binary data representation

“Binary” means “base 2”.

Digital computers represent all numbers using base-2 rather than base 10. For example:

42 in base 10: $4 \times \underline{10^1} + 2 \times \underline{10^0}$

42 in base 2:

$$1 \times \underline{2^5} + 0 \times 2^4 + 1 \times \underline{2^3} + 0 \times 2^2 + 1 \times \underline{2^1} + 0 \times 2^0 \\ = 1 \times 32 + 0 \times 16 + 1 \times 8 + 0 \times 4 + 1 \times 2 + 0 \times 1$$

42 as a sequence of binary digits (“bits”) 101010

All data values are represented in binary (base-2) at the machine level. “Bitwise” operators allow you to work directly with binary values.

2. What character represents the bitwise XOR operation? How does it differ from the OR operation?

Bitwise OR: combine two binary values by computing the result of the OR operation on each pair of binary digits.

Operator: |

Logic: 0|0=0, 0|1=1, 1|0=1, 1|1=1

Bitwise XOR: combine two binary values by computing the result of the XOR (exclusive or) operation on each pair of binary digits.

Operator: ^

Logic: 0^0=0, 0^1=1, 1^0=1, 1^1=0

Bitwise AND

Bitwise AND: combine two binary values by computing the result of the AND operation on each pair of binary digits.

Operator: `&`

Logic: $0 \& 0 = 0$, $0 \& 1 = 0$, $1 \& 0 = 0$, $1 \& 1 = 1$

3. What happens if you apply the bitwise operation on an integer value? (extra: what if we apply to floats)

The two-operand bitwise operators (`|`, `^`, `&`) perform a logical operation (OR, XOR, AND) on each pair of bits in the two operands.

The operands must be values belonging to an “integral” (integer-like) type.

E.g., `int`, `unsigned`, `long`, `unsigned long`, `char`, `unsigned char`, etc.

Bitwise operations may *not* be performed on floating point (`float` or `double`) values.

4. What is the result of $(15 \gg 2) \parallel 7$?
Handwritten notes: 3 above >>, 11 above ||, 7 above 7, right shift below >>, logical below ||

15 in base-2 is 1111

">>" is the right shift operator, shifting 1111 two bits to the right yields 0011, which is equal to 3.

Any non-zero integer is considered TRUE, so 3 is true.

If the left operand of \parallel is true, the entire expression is true (and the right operand is not evaluated.) All logical and relational operators yield 0 when false and 1 when true.

So, the result of $(15 \gg 2) \parallel 7$ is 1.

5. What is the result of $(15 \gg 2) | 7$?

$|$ is the bitwise OR operator

15 in binary is 1111

$(15 \gg 2)$ is 0011

7 in binary is 0111



```
  0011  
| 0111  
-----  
  0111
```

-- bit is 1 where either operand has a 1 bit
which is equal to 7

Exercise 14

- arrays and strings
- bitwise operations
- Breakout rooms 1–10 are “social”
- Use Slack to let us know if you have any questions!

Notes

Notes

Notes

Notes

Notes